

Covalent.

Workflow Orchestration for heterogeneous (HPC/Quantum) computing

Jack S. Baker, PhD
Quantum Algorithms
Researcher



Venkat Bala, PhD
HPC Engineer
Agnostiq





Overview

1. Introduction to Covalent
2. Covalent 101
3. Code Examples
 - a. Multiple dispatches, sub-lattices, results
 - b. Covalent Executors
 - c. C/C++ integration (Classical HPC)
4. Introduction to Quantum Machine Learning
5. Focus on Similarity Learning
6. Heterogeneous Similarity learning with Covalent (**hands-on workshop**)



Workflow Orchestration



Why is workflow orchestration is needed?

Primary Challenges:

- Explore new research problems
- Apply interdisciplinary techniques
- Push the boundaries of hardware
- Find hidden patterns in data

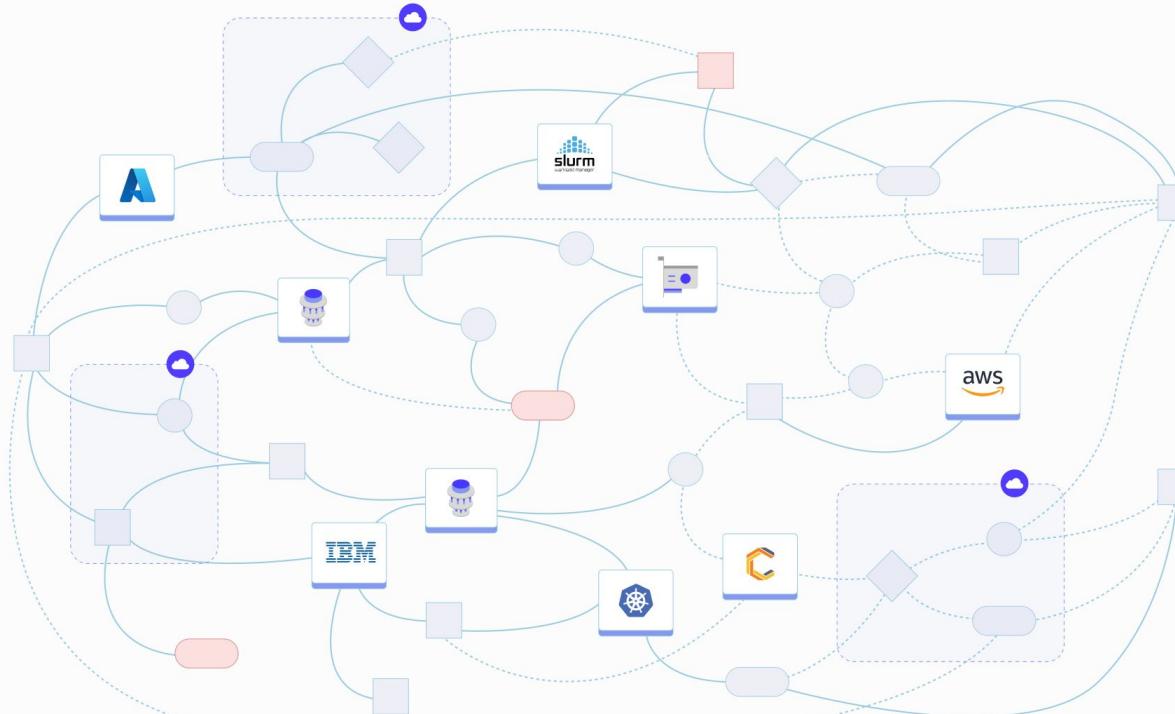
Additional Operational Challenges:

- Organizing and versioning experiments
- Managing and sharing data
- Managing accounts on multiple clusters
- Checkpointing long simulations
- Reproducing (your own) work
- Package version conflicts
- Compiling colleagues' code
- Incorporating legacy code
- Identifying optimal hardware resources
- Infrastructure management
- Working within a budget constraint

Lots of work, little reward... everyone comes up with their own bespoke solutions to these



Modern computational workflows are complex.



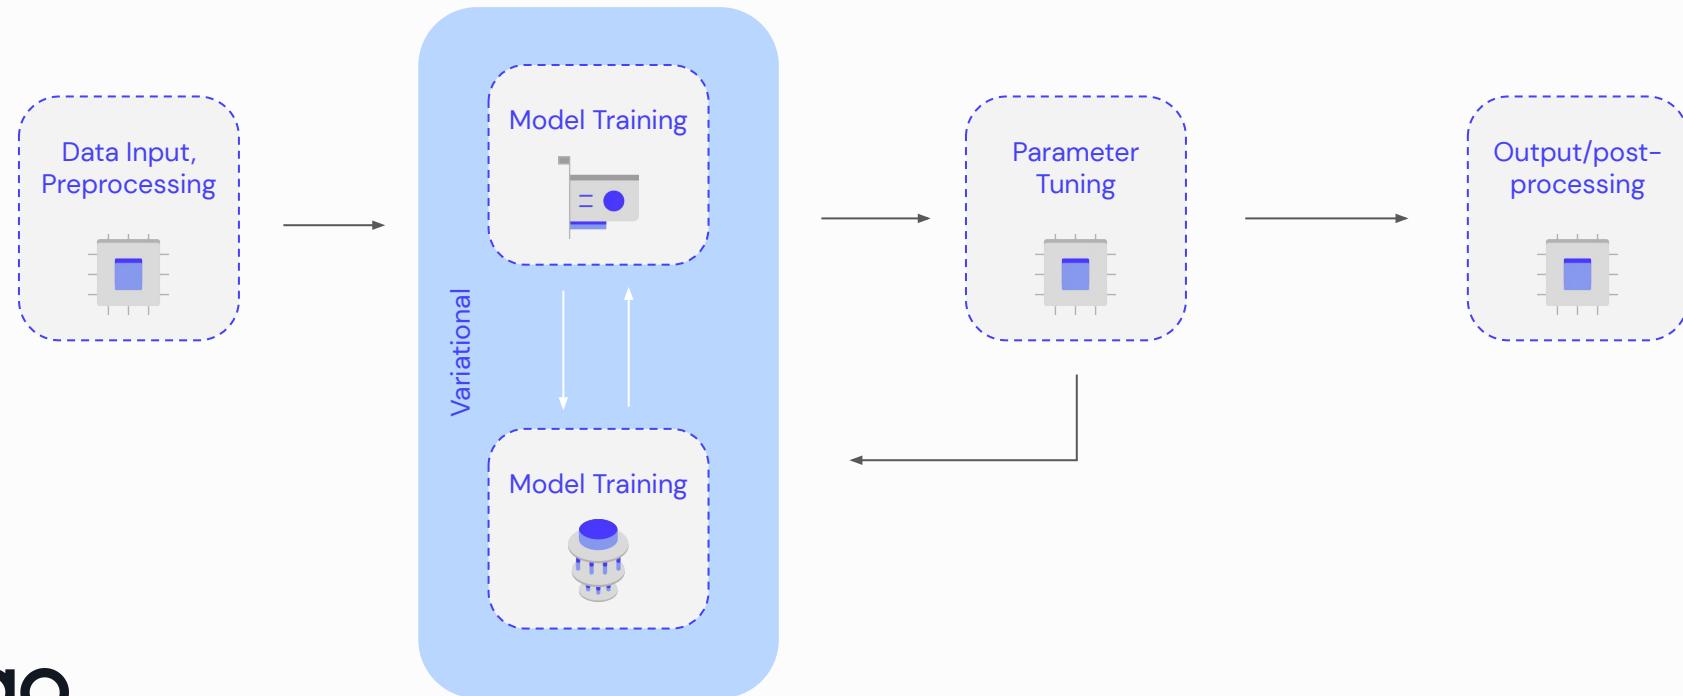
aq

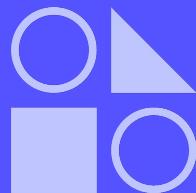
covalent.xyz



Quantum is driving the need for more heterogeneity

A simplified quantum machine learning workflow:





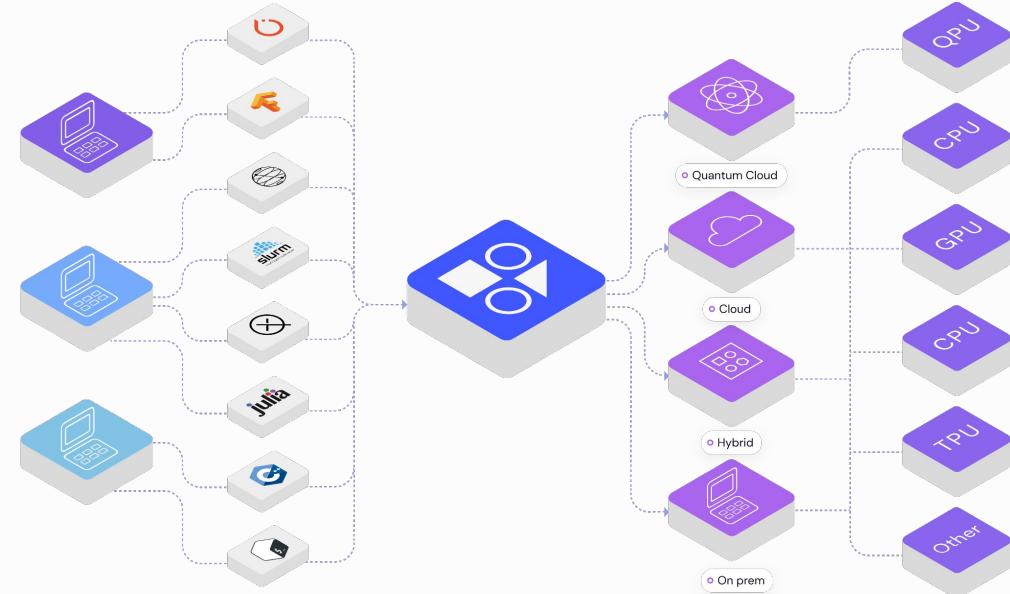
Covalent.



Covalent is an open source, free workflow orchestration platform for heterogeneous computing

Covalent.

- Support for quantum + HPC
- Distributed computation
- Heterogeneous hardware and software
- Workflow management interface



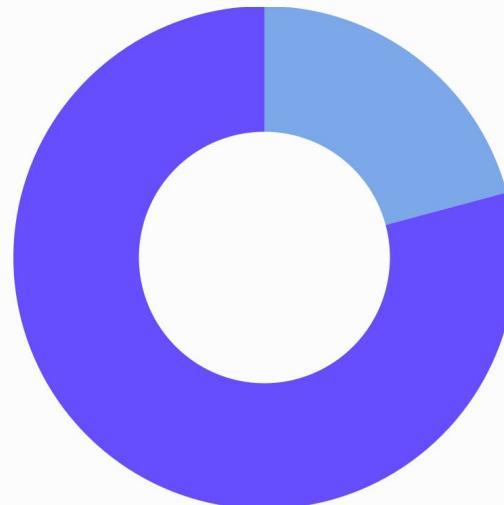
AgnostiqHQ/covalent

covalent.xyz



Why does Covalent exist?

Computational research



Time spent on
Writing,
Debugging,
Executing,
Maintaining,
Research code

Time spent on research



How Covalent can help



Solution 1.

Hybrid research / experiments. A single QML experiment now contains CPU + GPU/TPU + QPU.



Solution 2.

HPC research requires rapid prototyping and experimentation across software parameters.



Solution 3.

Execute high-throughput calculations. Run massively parallel jobs at scale across clusters.



Solution 4.

Avoid long HPC queue times and redeploy to different quantum queues using serverless HPC.





Where Covalent fits in the (Python) stack





Covalent supports a growing ecosystem of hardware and software

Classical Resources

- Slurm executor
- AWS Fargate executor
- AWS Batch executor
- Azure executor
- GCP executor
- Kubernetes executor



Quantum Resources

- AWS-Braket Jobs
- Qiskit runtime
(coming soon)
- More to come
- **Write your own plugin!**



Software Packages

Any and all packages!



Languages



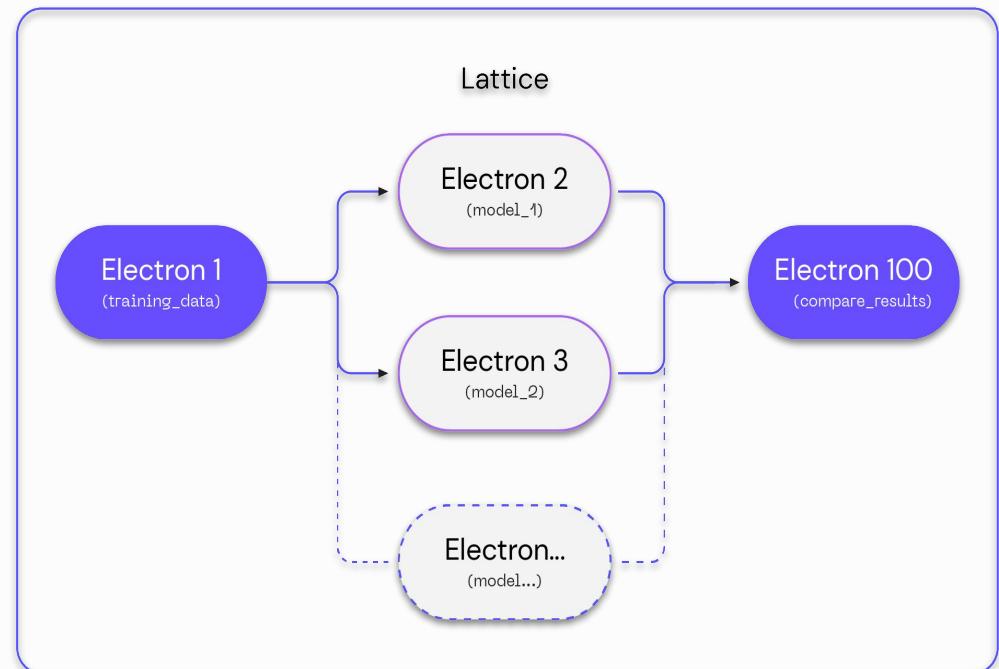


Covalent 101: The Basics



Covalent Concepts

- Workflows are comprised of inter-dependent tasks
- Workflows are called “lattices”
- Tasks are called “electrons” – these are fundamental blocks
- Tasks can be written in other languages
- Tasks can each run on different hardware backends





Minimal code changes required

ooo

```
1 def task(x, y):
2     z = optimize(x, y)
3     return z
4
5 def train(model):
6     import tensorflow as tf
7     tf.train(model)
8     return model
9
10 def workflow(inputs):
11     model = task(inputs[0], inputs[1])
12     return train(model)
13
14 print(workflow([1,2]))
```

ooo



ooo

```
1 import covalent as ct
2
3 @ct.electron(executor=AWSBatchExecutor())
4 def task(x, y):
5     z = optimize(x, y)
6     return z
7
8 @ct.electron(executor=SlurmExecutor())
9 def train(model):
10    import tensorflow as tf
11    tf.train(model)
12    return model
13
14 @ct.lattice
16 def workflow(inputs):
17    model = task(inputs[0], inputs[1])
18    return train(model)
19
20 dispatch_id = ct.dispatch(workflow)([1,2])
21 print(ct.get_result(dispatch_id))
```



Covalent UI

Completed | 4 / 4

```
graph LR; 1((1)) -- x --> task((task)); 2((2)) -- y --> task; task -- model --> train((train))
```

Overview

Started - Ended
Sep 19, 18:49:28 - Sep 19, 18:49:29

Runtime
< 1min

Directory
/home/venkat/tu...8d-a34b-737618ac5586

Input
{"args": ["[1, 2]"], "kwargs": {}}

Result
1

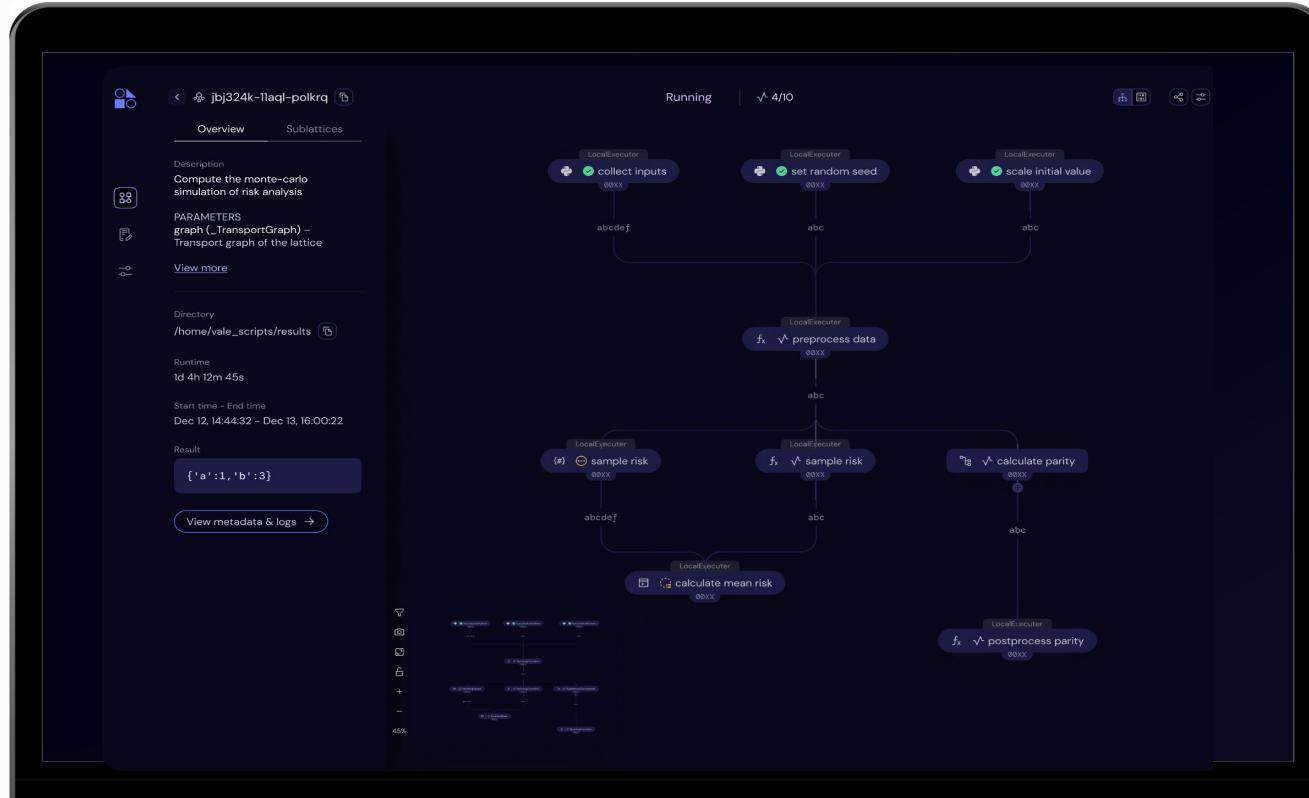
Executor: dask

```
@ct.lattice
def workflow(inputs):
    model = task(inputs[0], inputs[1])
    return train(model)
```

A sidebar on the left contains icons for file, folder, search, and other navigation functions.



Visualize and monitor complex workflows





Inspect task-level results and metadata

jbj324k-1laql-polkrq

Running | ✓ 4/10

Overview Sublattices

Description: Compute the monte-carlo simulation of risk analysis

PARAMETERS: graph (.TransportGraph) - transport graph of the lattice

View more

Directory: /home/vale_scripts/results

Runtime: 1d 4h 12m 45s

Start time - End time: Dec 12, 14:44:32 - Dec 13, 16:00:22

Result: { 'a':1, 'b':3}

View metadata & logs →

LocalExecutor collect inputs (0xx) abcdef

LocalExecutor set random seed (0xx) abc

fx ✓ preprocess data (0xx) abc

[x] ⊕ sample risk (0xx) abcdef

fx ✓ sample risk (0xx) abc

LocalExecutor calculate mean risk (0xx)

sample risk

Status: Running

Description: Samples risk value for various points according to a given probability distribution

Start time - End time: Dec 12, 14:44:32 - Dec 13, 16:00:22

```
# @ct.electron
def sample_risk (x, y) :
    ...
    ...
```

Runtime: 12m 45s

Executor: Local conda_env ...

Result: [0.1, 0.2, 0.1, ..., 0.01, 0.3]

View metadata & logs →



Organize jobs across projects and experiments

Projects

White Silver Doberman 3 days 8/8 ✓ 3h 20m 15/20

Gray Titanium Beagle 6 days 2/5 ✓ 14h 49m 9/17

Yellow Lead Great Dane 2/5 ✓ 14h 49m 9/17

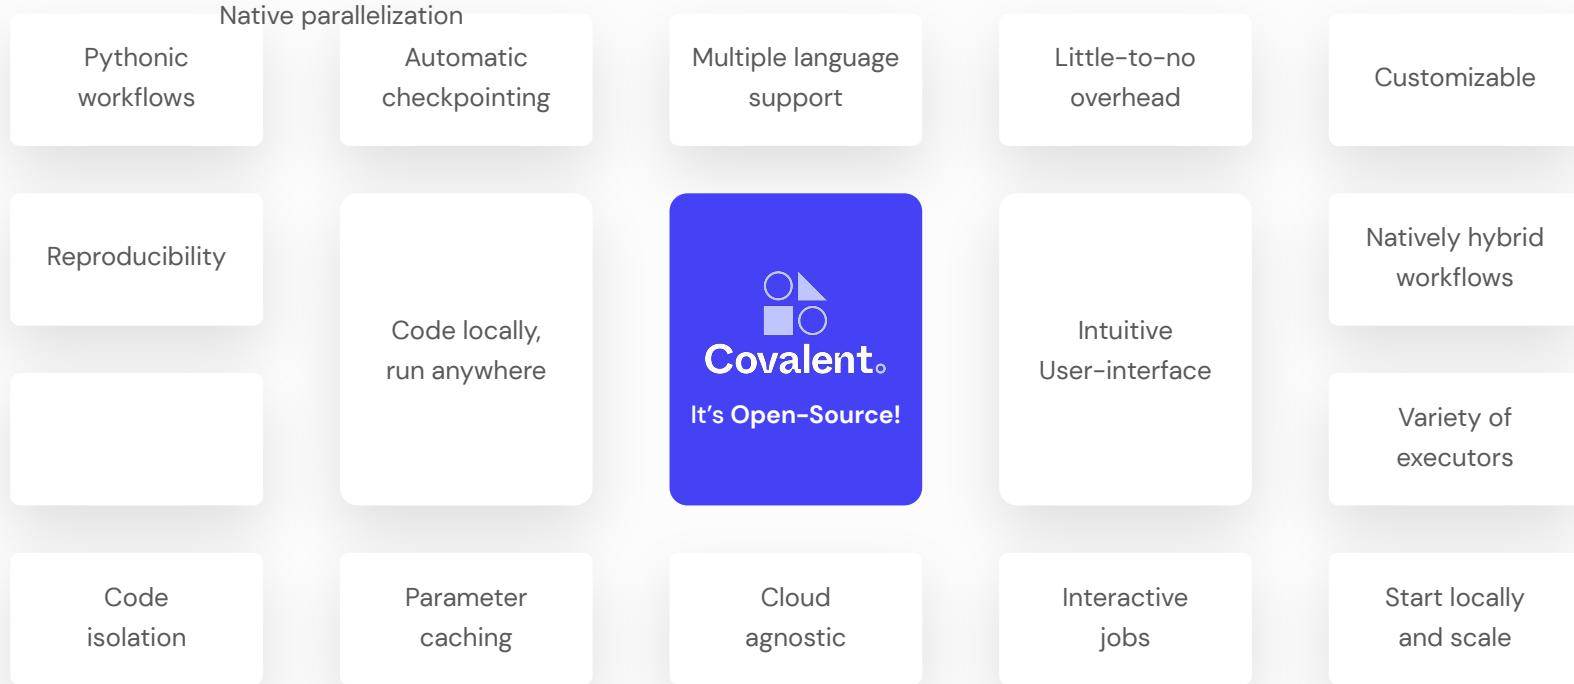
Red Aluminium Labrador 2/5 ✓ 14h 49m 9/17

All Experiments Dispatches Show Add new Search

Title	Status	Tags	Runtime	Start time	Last updated
Steel Pelican Hooper 5					
Platinum Heron Chimera 6					
Navy Bronze Beagle 22/22	react rust node js	3 days	12 Jan, 04:25:00	13 Jan, 08:31:00	
Gray Silver Newfoundland 69/79	react rust node js	14 days	12 Jan, 03:07:00	13 Jan, 08:12:00	
Teal Copper Spitz 20/20	react rust node js	13h 36m	12 Jan, 01:08:00	13 Jan, 08:04:00	
Green Steel Dalmatian 13/13	react rust node js	2 days	12 Jan, 00:55:00	13 Jan, 07:53:00	
Maroon Rattlesnake Wyvern 17/25	react python rust c2	4 days	08 Jan, 23:42:00	07 Jan, 00:53:00	



You can contribute too!





Code Samples

Example 1: Curve Fitting

SETUP



- Clone the Github repository

```
git clone https://github.com/AgnostiqHQ/tutorials_coalent_ieee_2022.git
```

- Download [Miniconda](#)
- Install the conda environment (`environment.yml` in root directory)

```
conda env create -f environment.yml
```

- Activate environment

```
conda activate ieee_coalent
```

- Add to kernel list

```
python -m ipykernel install --user --name=ieee_coalent
```



covalent start

```
covalent start --ignore-migrations
```



1. Functionalize your code

```
import numpy as np
import matplotlib.pyplot as plt
x = np.arange(10)
Y = np.random.random(10)

# Fit
fit = np.polyfit(x, y, 3)
xnew = np.linspace(x[0], x[-1], 50)
y_new = fit(xnew)

# plot
fig, ax = plt.subplots()
plt.plot(x, y, 'o', xnew, y_new)
plt.xlim([x[0]-1, x[-1] + 1])
```

```
import numpy as np
import matplotlib.pyplot as plt

def fit_xy(x,y,order):
    z = np.polyfit(x,y,order)
    return np.poly1d(z)

def plot_fit(x,y,xnew,fit):
    y_new=fit(xnew)
    fig,ax=plt.subplots()
    plt.plot(x,y, 'o', x_new, y_new)
    plt.xlim([x[0]-1, x[-1] + 1])
    return fig

def cfit(x,y,order):
    x_new = np.linspace(x[0], x[-1], 50)
    fit=fit_xy(x,y,order)
    return plot_fit(x=x,y=y,xnew=x_new,fit=fit)
```



2. Decorate your functions

```
def fit_xy(x,y, order):
    z = np.polyfit(x, y, order)
    return np.poly1d(z)
```



```
import covalent as ct
@ct.electron
def fit_xy(x,y, order):
    z = np.polyfit(x, y, order)
    return np.poly1d(z)
```

```
def plot_fit(x,y,xnew,fit):
    y_new=fit(xnew)
    fig,ax=plt.subplots()
    plt.plot(x,y,'o', x_new, y_new)
    plt.xlim([x[0]-1, x[-1] + 1 ])
    return fig
```



```
@ct.electron
def plot_fit(x,y,xnew,fit):
    y_new=fit(xnew)
    fig,ax=plt.subplots()
    plt.plot(x,y,'o', x_new, y_new)
    plt.xlim([x[0]-1, x[-1] + 1 ])
    return fig
```



3. Create your workflow

```
def cfit(x, y, order):
    x_new = np.linspace(x[0], x[-1], 50)
    fit = fit_xy(x, y, order)
    return plot_fit(x, y, x_new, fit)
```

```
@ct.lattice
def cfit(x, y, order):
    x_new = np.linspace(x[0], x[-1], 50)
    fit = fit_xy(x=x, y=y, order=order)
    return plot_fit(x=x, y=y, xnew=x_new, fit=fit)
```



4. Dispatch your workflow

```
x = np.arange(10)  
y = np.random.random(10)  
dispatch_id = ct.dispatch(cfit)(x, y, 10)
```

5. Get the result

```
result = ct.get_result(dispatch_id, wait=True)
```



6. View results on UI

The screenshot shows the Covalent UI interface. On the left, there's a detailed view of a workflow run with the ID `a454b34a...-de05229efa96`. The "Overview" section includes:

- Started - Ended:** Sep 20, 18:45:01 - Sep 20, 18:45:01
- Runtime:** < 1sec
- Directory:** /home/venkat/tu...59-bf92-de05229efa96
- Input:** `{"args": [], "kwargs": {"x": "[0 1 2 3 4 5 6 7 8 9]”}}`
- Result:** `AxesSubplot(0.111372, 0.131829; 0.851421x0.78831)`
- Executor:** `task`

On the right, the status bar indicates "Completed" with "9 / 9". Below it, a workflow graph is displayed:

```
graph TD; fit_xy((fit_xy)) --> fit((fit)); fit --> plot_fit((plot_fit))
```

The nodes in the graph are:

- `fit_xy` (highlighted with a green oval)
- `fit` (highlighted with a blue oval)
- `plot_fit` (highlighted with a green oval)

The graph shows a sequential flow from `fit_xy` to `fit`, and then from `fit` to `plot_fit`.



Internals

- Decorators declare intent
- Directed Acyclic Graph (DAG) is built by sequentially inspecting the electrons (TransportGraph)
- Internally **Covalent** uses a local [Dask](#) cluster to execute the workflow
- Default cluster configuration can be set at covalent start time
 - Default workers = # of cores
 - Memory per workers = Total memory/Number of workers



Multiple Dispatches



Multiple dispatches

- Repeated experiments
- Large parameter sweeps, hyperparameter optimization in machine learning
- Dispatch a lattice with different parameters each time

```
inputs = [{"x": np.arange(N), "y": np.random.random(N), "order": np.random.randint(5, N)} for _ in range(5)]
dispatches = [ct.dispatch(exp)(input['x'], input['y'], input['order']) for input in inputs]
results = [ct.get_result(dispatch_id, wait=True) for dispatch_id in dispatches]
```

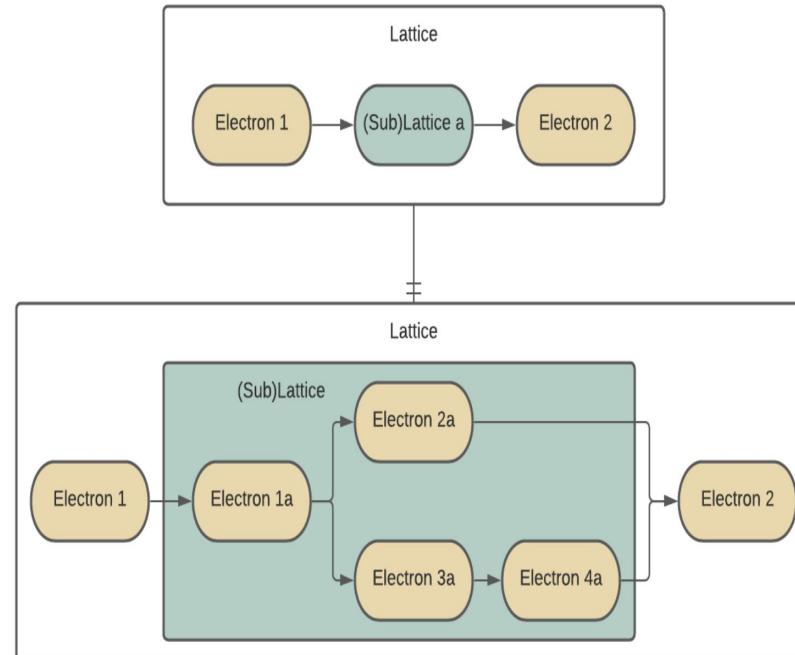


Sublattices



Sublattices

- A **lattice** converted into an **electron**
- Transforms a workflow into a node in a much bigger workflow
- Composition
- Makes creating larger workflows from existing ones much easier
- Add an electron **decorator** to an existing lattice





Motivation

- Nested experiments
- Created an experiment (lattice) consisting of several tasks (electrons)
- Dispatch the experiment to Covalent for execution with a specific set of parameters
- Create a new workflow that takes in a range of parameters as input
- Invoke the earlier workflow as a node
- Can recursively construct very large nested workflows

```
@ct.electron
def task1(*args, **kwargs):
    ...
    return result

@ct.electron
def task2(*args, **kwargs):
    ...
    return result

@ct.lattice
def experiment(**params):
    r1 = task1(*args, **kwargs)
    r2 = task2(*args, **kwargs)
    return r1 + r2

# Create sub-lattice
run_experiment = ct.electron(experiment)

@ct.lattice
def experiment_suite(**parameters):
    for param in parameters
        run_experiment(**param)
    return
```

Sublattice



Overview

Started - Ended
Sep 20, 15:11:06 - Sep 20, 15:11:31

Runtime
< 1min

Directory
`/home/venkat/tu...c2-a862-40bf8f2f74ed`

Input

```
{"args": [{"x": array([0, 1, 2, 3, 4, 5, 6, 7, 8,
```

Result

None

Executor: **task**

```
@ct.lattice
def multiple_curve_fits(inputs: List):
    all_results = []
    for input in inputs:
        result = cfit_sublattice(x=input["x"], y = i
    all_results.append(result)
```

Completed 200 / 200

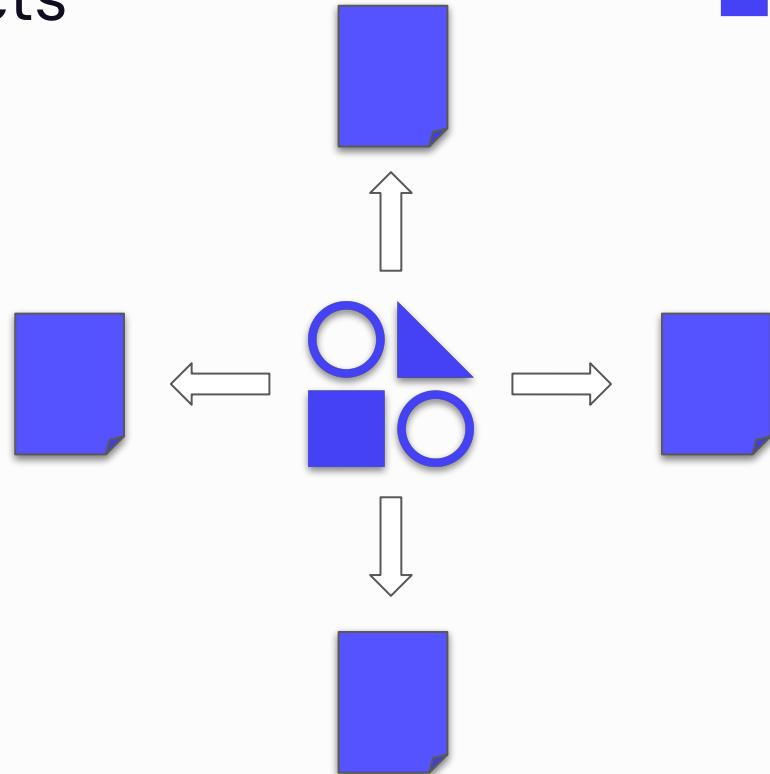


Metadata



Results Manager/Result Objects

- Storing, updating and retrieving workflow results
- Decoupled from dispatch server
- Users can query and track workflow progress without needing source code
- Experiment outcomes can be safely stored
- Build in checkpointing
- Easy API to retrieve result





Code Samples

Example 2: Eigenvalues



Matrix Eigenvalues

- Important in classical/quantum mechanics
- Can be computationally heavy
- Demonstrate concept of **remote executors** in Covalent
- Can display beautiful symmetries!



Matrix Eigenvalues

```
import numpy as np

np.random.choice([-1, 0, 1], 25).reshape(5, 5)

array([[ 0,  1,  0,  0,  1],
       [ 1,  0,  1,  0, -1],
       [ 0,  0,  0, -1, -1],
       [ 1, -1,  0,  1, -1],
       [-1,  1, -1, -1,  0]])
```

Compute the eigenvalues of a lot of such matrices and view the results



Workflow Electrons

```
@ct.electron
def generate_random_matrix(integer_set, nrows, ncols):
    return np.random.choice(integer_set, nrows*ncols).reshape(nrows, ncols)

@ct.electron
def compute_eigenvalues(matrix):
    eig, _ = np.linalg.eig(matrix)
    return eig

@ct.electron
def get_real_part(eigenvalues):
    return [np.real(eig) for eig in eigenvalues]

@ct.electron
def get_imag_part(eigenvalues):
    return [np.imag(eig) for eig in eigenvalues]
```



Lattice

```
@ct.lattice
def eigenvalue_workflow(N: int, batch_size: int):
    real_part = []
    imag_part = []
    for index in range(batch_size):
        matrix = generate_random_matrix(N)
        eigenvalues = compute_eigenvalues(matrix)
        real_part.append(get_real_part(eigenvalues))
        imag_part.append(get_imag_part(eigenvalues))
    return np.asarray(real_part).flatten(), np.asarray(imag_part).flatten()
```

Workflow Graph



ce66051f...-c8e270000e9e

Overview

Started - Ended
Sep 21, 17:52:39 - Sep 21, 17:52:43

Runtime
< 1min

Directory
/home/venkat/tu...c8-a251-c8e270000e9e

Input

```
{"args": ["5", "20"], "kwargs": {}}
```

Result

```
(array([-1.94457628e+00, -1.94457628e+00, 6.7076084
```

Executor: **task**

```
@ct.lattice
def eigenvalue_workflow(N: int, batch_size: int):
    real_part = []
    imag_part = []
    for index in range(batch_size):
        matrix = generate_random_matrix(N)
        eigenvalues = compute_eigenvalues(matrix)
```

Completed | 100 / 100



Executors



Executors

- With Covalent we can choose to execute certain tasks of a workflow on different hardware/machines of our choice
- This is made possible through the use of **executors**
- Executors allow the user to dictate the execution environment/platform of a particular task in the workflow
- Executors are plugins to Covalent and can be installed separately via **pip**
- Supported executors
 - Local, Dask, SSH, SLURM, AWS (EKS, ECS, Batch, Lambda, Braket), Qiskit Runtime (in development)



SSHExecutor

```
pip install covalent-ssh-plugin
```



SSHExecutor

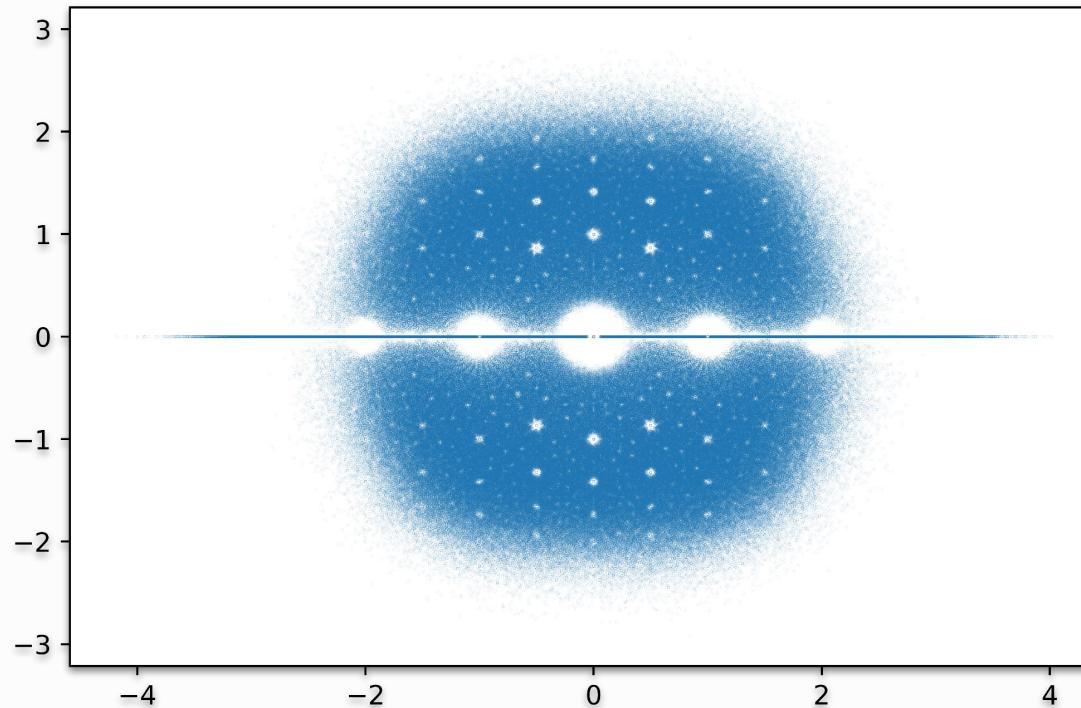
```
from covalent.executor import SSHExecutor

ssh = SSHExecutor(username="<username>", hostname="<hostname>",
                  ssh_key_file="<ssh private key>")

@ct.electron(executor=ssh)
def compute_heavy_task(*args, **kwargs):
    ...
    return result

@ct.lattice
def workflow(*args, **kwargs):
    ...
    res = compute_heavy_task(*args, **kwargs)
    ...
    return res
```

Result





SLURMExecutor

```
pip install covalent-slurm-plugin
```

SLURM Executor



```
from covalent.executor import SlurmExecutor

slurm = SlurmExecutor(username="xyz", address="cluster.example.edu",
                      ssh_key_file="~/.ssh/ssh_private_key", options={"ntasks": 1, "cpus-per-task": 5,
                      "partition": "compute", "nodelist": "node001"})

@ct.electron(executor=slurm)
def compute_heavy_task(*args, **kwargs):
    ...
    return result

@ct.lattice
def workflow(*args, **kwargs):
    res = compute_heavy_task(*args, **kwargs)
    return res
```



AWS Executor Plugins

AWS Executor Plugins

- Covalent supports a wide range of AWS specific cloud executors
- These executors allow users to dispatch electrons to different AWS compute services
- Supported AWS services
 - AWS Batch
 - AWS Lambda
 - AWS EC2
 - AWS Braket
 - AWS ECS/Fargate
 - AWS EKS





AWS Executors

- AWS Braket
 - Suitable for hybrid workflows (mix of classical and quantum resources)
- AWS Lambda
 - Embarrassingly parallel workflows, several short lived electrons (light weight)
- AWS ECS
 - Useful for moderate to heavy workloads (low memory requirements)
- AWS Batch
 - Useful for heavy compute workloads (high CPU/memory)
- AWS EC2
 - General purpose compute workloads, can be tuned to workflow's compute requirements



AWS Lambda Executor

```
import covalent as ct
from covalent.executor import AWSLambdaExecutor

awslambda = AWSLambdaExecutor(credentials="~/.aws/credentials",
                               profile="default",
                               region="us-east-1",
                               lambda_role_name="CovalentLambdaExecutionRole",
                               s3_bucket_name="covalent-lambda-job-resources")

@ct.electron(executor=awslambda)
def task(*args, **kwargs):
    ...
    return result

@ct.lattice
def workflow(*args, **kwargs):
    res = task(*args, **kwargs)
    return res
```



Fully configurable!!

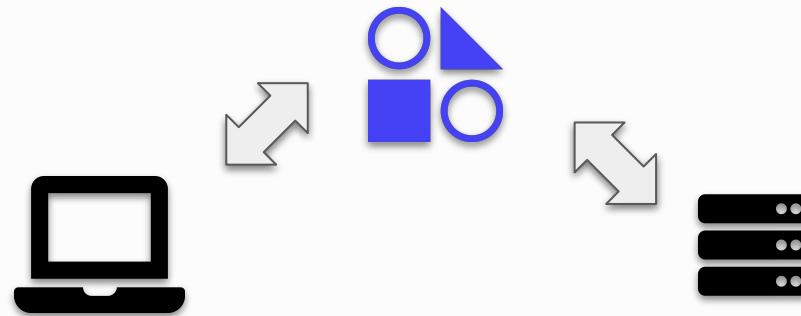


Extra features



File transfers

- Transfer file from local to remote machine (vice-versa)
- AWS S3 buckets supported
- pre/post execution to upload/download files
- Specified as electron decorator arguments





File Transfer API

```
import covalent as ct

@ct.electron(
    files=[ct.fs.FileTransfer('/home/ubuntu/src_file', '/home/ubuntu/dest_file')]
)
def my_task(files=[]):
    from_file, to_file = files[0]
    with open('/home/ubuntu/dest_file', 'r') as f:
        return f.read()

@ct.lattice
def file_transfer_workflow():
    return my_task()

# Dispatch the workflow
dispatch_id = ct.dispatch(file_transfer_workflow)()
```



Electron dependencies (Deps)

- `Deps` provide a way to specify the electron level dependencies in Covalent
- `DepsPip`
 - Install any Python libraries required by an `electron`
- `DepsBash`
 - Execute any `Bash` command/script before/after electron execution
- `DepsCall`
 - Invoke any Python function before/after execution of an electron



Electron dependencies

```
import covalent as ct
from covalent import DepsPip, DepsBash, DepsCall

def execute_before_electron(a, b):
    ...

def shutdown_after_electron():
    ...

@ct.electron(
    deps_pip=DepsPip(packages=[ "numpy==0.23", "qiskit" ]),
    deps_bash=DepsBash(commands=[ "echo $PATH", "ssh foo@bar.com" ]),
    call_before=DepsCall(execute_before_electron, args=(1, 2)),
    call_after=DepsCall(shutdown_after_electron),
    executor=slurm
)
def task(*args, **kwargs):
    ...
```

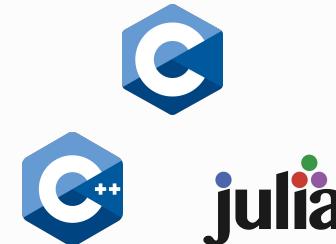


C/C++ support



Interfacing with other languages

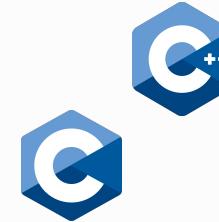
- Covalent is designed to be **language agnostic**
- If there are bindings between Python and the language of choice, workflows can be constructed where electrons invoke code written in different languages
- Here we demonstrate how one can interface with a module written in C/C++ through Covalent
- More language support coming
 - Julia
 - Fortran
 - (Request your favorite language on GitHub)





Why is it important?

- Interfacing with legacy code
- Quantum simulations on classical hardware expensive (# of qubits)
- A lot of quantum simulators are available in C/C++
 - [NVIDIA cuQuantum, QODA](#)
 - GPU acceleration for quantum simulations
 - [Intel Quantum Simulator \(IQS\)](#) (C/C++)
 - Accelerate quantum simulations on Intel CPUs
 - [Quantum++](#) (C++ 11)
 - [Qrack](#) (C++/OpenCL)





Code Samples

Example 3: Estimate PI in C

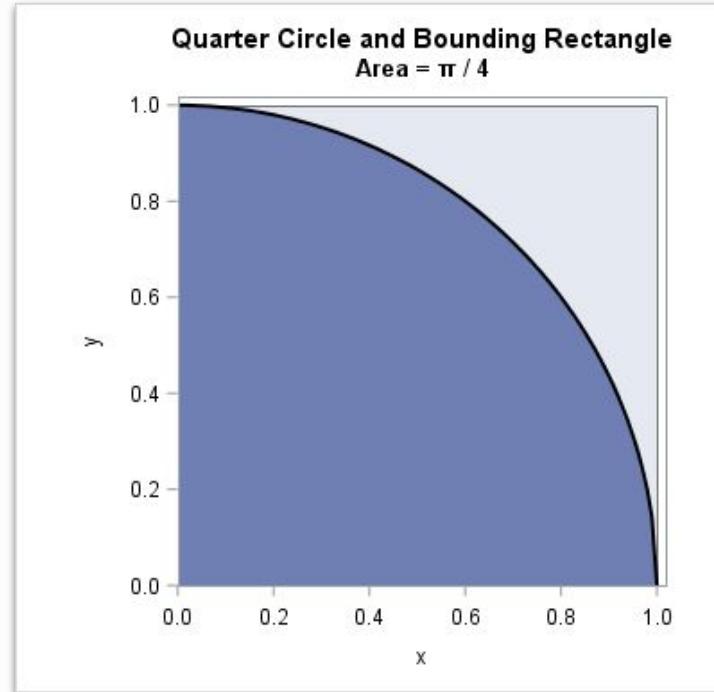


Interfacing with C

- CPython (most widely used) is written purely in C
- Python module written in C
- Compile the module to a shared library
- Use the functions defined in the module in **Covalent** workflows
- [Python/C API](#)
- Wrappers like [SWIG](#) can also be leveraged



Estimate PI: Area under the curve



Python/C API



```
#include <Python.h>

static PyObject* compute_pi(PyObject* self, PyObject* args) {
    unsigned int partitions;
    double area = 0.0;
    double dh;

    if (!PyArg_ParseTuple(args, "I", &partitions)) {
        return NULL;
    }
    dh = 1.0/partitions;

    for (int i = 0; i < partitions; i++) {
        double x = dh * (i - 0.5);
        area += (4.0/(1.0 + x*x));
    }
    // pi approximation
    return PyFloat_FromDouble(area*dh);
}
```



Create Python module

```
static PyMethodDef PiapproxMethods[] = {
    {"compute_pi", compute_pi, METH_VARARGS, "Compute an approximation to PI"},
    {NULL, NULL, 0, NULL}           /* Sentinel */
};

static struct PyModuleDef cpiapprox = {
    PyModuleDef_HEAD_INIT,
    "cpiapprox",      /* name of module */
    NULL,             /* module documentation, may be NULL */
    -1,               /* size of per-interpreter state of the module,
                      or -1 if the module keeps state in global variables. */
    PiapproxMethods
};

PyMODINIT_FUNC PyInit_cpiapprox(void) {
    return PyModule_Create(&cpiapprox);
}
```



Build and Install

- setup.py

```
from distutils.core import setup, Extension
cpiapprox = Extension('cpiapprox', sources=['main.c'],
                      extra_compile_args=["-O3"])
setup(name='cpiapprox', version='1.0',
      description='A approximation to PI', ext_modules=[cpiapprox])
```

- Build/install the module

```
python setup.py install
```



Covalent workflow

```
import covalent as ct
import cpiapprox

@ct.electron
def compute_pi(n):
    return cpiapprox.compute_pi(n)

@ct.lattice
def workflow():
    res = []
    for i in [100, 1000, 100000]:
        res.append(compute_pi(i))
    return res

dispatch_id = ct.dispatch(workflow)()
res = ct.get_result(dispatch_id, wait=True)

print(res.result)

[3.161499736951266,
 3.1435917356731236,
 3.1416126534981603]
```

The screenshot shows the Covalent UI interface for a completed workflow run. The main panel displays the workflow details:

- Status:** Completed
- Progress:** 3/3
- Started - Ended:** Jul 06, 02:48:10 - Jul 06, 02:48:10
- Runtime:** 00:00:00
- Os:** Linux
- Directory:** /home/venkat/tu..pi_2022/notebooks/results
- Result:** [3.161499736951266, 3.1435917356731236, 3.1416126534981603]
- Executor:** task
- cache_dir:** /home/venkat/.cache/covalent
- current_env_on_conda_fail:** False

On the right side, there are several small preview cards for the tasks: three cards labeled "compute_pi" and one card labeled "compute_pi".



Code Samples

Example 4: Vector
Arithmetic in C++



Interfacing with C++

- Pybind11 is a lightweight header-only library that exposes C++ types in Python and vice-versa
- It supports a lot of C++ core features
 - Functions (passing by value and reference)
 - Instance/static method
 - Function overloading
 - Templates
 - STL types
 - Iterators, ranges
 - single/multiple inheritance



C++ with Covalent

- As a very minimal example of, we write a simple vector arithmetic library in C++ and expose it as a Python package to Covalent via `pybind11`
- We use core C++ STL types (`std::vector`) and STL algorithms to implement three functions
 - `Vecadd`
 - Add two `std::vector` of type double (elementwise)
 - `Vecmul`
 - Multiply two `std::vector` of type double (elementwise)
 - `Vecdiv`
 - Divide two `std::vector` of type double (elementwise)



Vector arithmetic

```
#include <vector>
#include <pybind11/pybind11>
#include <numeric>
#include <algorithm>
#include <functional>

using dvec = std::vector<double>;

dvec vecadd(const dvec& a, const dvec& b) {
    dvec c;
    std::transform(a.begin(), a.end(), b.begin(),
                  std::back_inserter(c), std::plus<double>());
    return c;
}
dvec vecmul(const dvec& a, const dvec& b) {
    dvec c;
    std::transform(a.begin(), a.end(), b.begin(),
                  std::back_inserter(c), std::multiplies<double>());
    return c;
}
dvec vecdiv(const dvec& a, const dvec& b) {
    dvec c;
    std::transform(a.begin(), a.end(), b.begin(),
                  std::back_inserter(c), std::divides<double>());
    return c;
}
```



Build Python interface

```
// main.cc
PYBIND11_MODULE(cpparthimetic, m) {
    m.def("vecadd", &vecadd, "Add two python lists");
    m.def("vecmul", &vecmul, "Multiply two python lists");
    m.def("vecdiv", &vecdiv, "Divide two python lists elementwise");
}
```



Compiling C++ to generate module

- setup.py

```
# setup.py
from setuptools import setup
from pybind11.setup_helpers import Pybind11Extension, build_ext

cpparthimetic_module = Pybind11Extension('cpparthimetic',
                                         sources=['main.cc'])

setup(name = 'cpparthimetic',
      version='0.1.0',
      author = 'venkat Bala',
      author_email="venkat@agnostiq.ai",
      cmdclass={"build_ext": build_ext},
      ext_modules=[cpparthimetic_module])
```

- Build/install module

```
python setup.py install
```



Covalent workflow

```
import covalent as ct
import random
import cpparithmetic

@ct.electron
def generate_list(N: int):
    return [random.uniform(0.1, 1) for i in range(N)]

@ct.electron
def cppadd(a: List[float], b: List[float]):
    return cpparithmetic.vecadd(a, b)

@ct.lattice
def cppadd_workflow(size):
    a = generate_lists(size)
    b = generate_lists(size)
    return cppadd(a, b)

dispatch_id = ct.dispatch(cppadd_workflow)(2**20)
add_result = ct.get_result(dispatch_id, wait=True)
```



C++ workflow

8480e650...-436290a1e5d6

Completed | 5 / 5

Started - Ended
Sep 21, 23:26:12 – Sep 21, 23:26:13

Runtime
< 1min

Directory
/home/venkat/tu...10-8ccd-436290a1e5d6

Input
{"args": ["16"], "kwargs": {}}

Result
[8.262604904861552, 5.7566025790075575, 11.225241943]

Executor: **task**

```
@ct.lattice
def cppadd_workflow(size):
    a = generate_lists(size)
    b = generate_lists(size)
    c = cppadd(a, b)
    return c
```

```
graph TD
    A((generate_lists)) --> C((cppadd))
    B((generate_lists)) --> C
```

The screenshot shows the Covalent UI for a completed C++ workflow. The workflow consists of three main components: two parallel tasks named 'generate_lists' and one final task named 'cppadd'. The 'generate_lists' tasks are completed, indicated by green checkmarks. The 'cppadd' task is also completed. The input to the workflow is a JSON object with 'args' and 'kwargs' fields. The output is a list of three floating-point numbers. The executor is set to 'task'. The code block at the bottom shows the Python code used to define the workflow, utilizing the Covalent library's '@ct.lattice' and 'cppadd_workflow' decorators.



Summary



Covalent

- Covalent is fully open source and under going active development
- Contributions, feature requests are welcomed
- Source code
 - <https://github.com/AgnostiqHQ/covalent>
- Documentation
 - <https://covalent.readthedocs.io/en/stable/>
- Upcoming features/issues/bug reports
 - <https://github.com/AgnostiqHQ/covalent/issues>
- Latest/pre-releases available on PyPI
 - <https://pypi.org/project/covalent/>





pip install covalent



A (very brief) Introduction to Quantum Machine Learning



Computer Science

Physics

Quantum
Computing

Math

Engineering

Key

Core discipline

aQ



Data Science

Machine
Learning

Materials

Computer Science

Finance

Physics

Cryptography

Quantum Computing

Mathematics

Engineering

Drug
Development

Key

Quantum
Chemistry

Biology

Core discipline

Application

aq



Data Science

Machine Learning

Materials

Computer Science

Finance

Physics

Cryptography

Quantum Computing

Mathematics

Engineering

Drug Development

Key

Quantum Chemistry

Biology

Core discipline

Application

aq



What is Quantum Machine Learning?



Before addressing Quantum Machine Learning (QML), we must first address ML



Model parameters are tuned while being fed a training dataset



The predictive capability of a model is assessed using a testing dataset



What is Quantum Machine Learning?



Before addressing Quantum Machine Learning (QML), we must first address ML



Model parameters are tuned while being fed a training dataset



The predictive capability of a model is assessed using a testing dataset



ML: The construction of predictive models without *explicit instruction*.



What is Quantum Machine Learning?



Before addressing Quantum Machine Learning (QML), we must first address ML



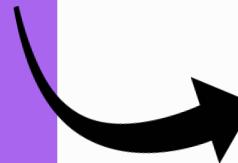
ML: The construction of predictive models without *explicit instruction*.



Model parameters are tuned while being fed a training dataset



The predictive capability of a model is assessed using a testing dataset

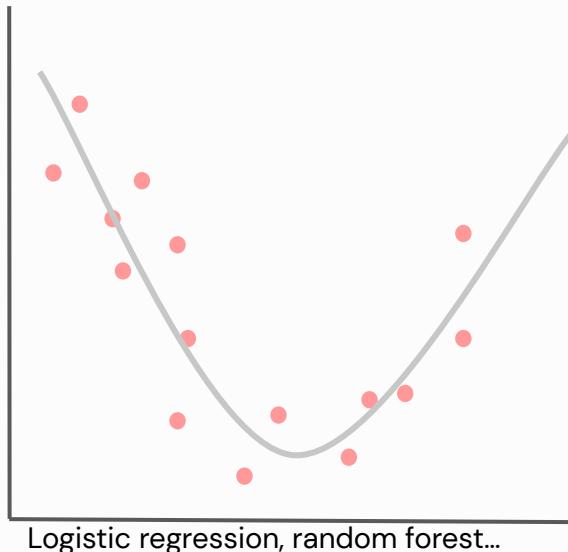


Often, validation is performed between training and testing



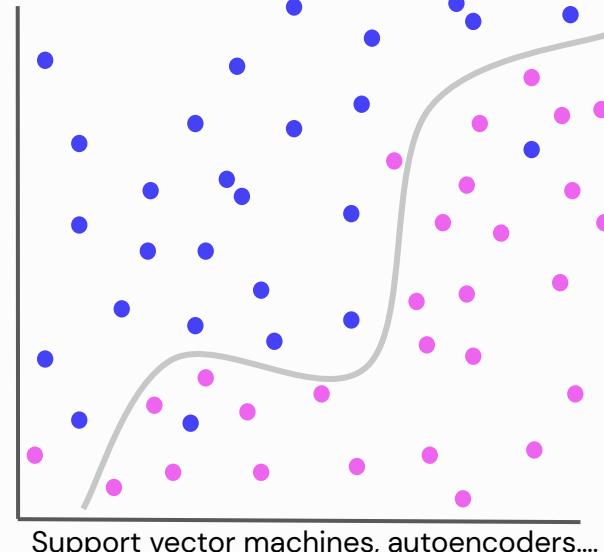
What is Quantum Machine Learning?

Regression



Logistic regression, random forest...

Classification



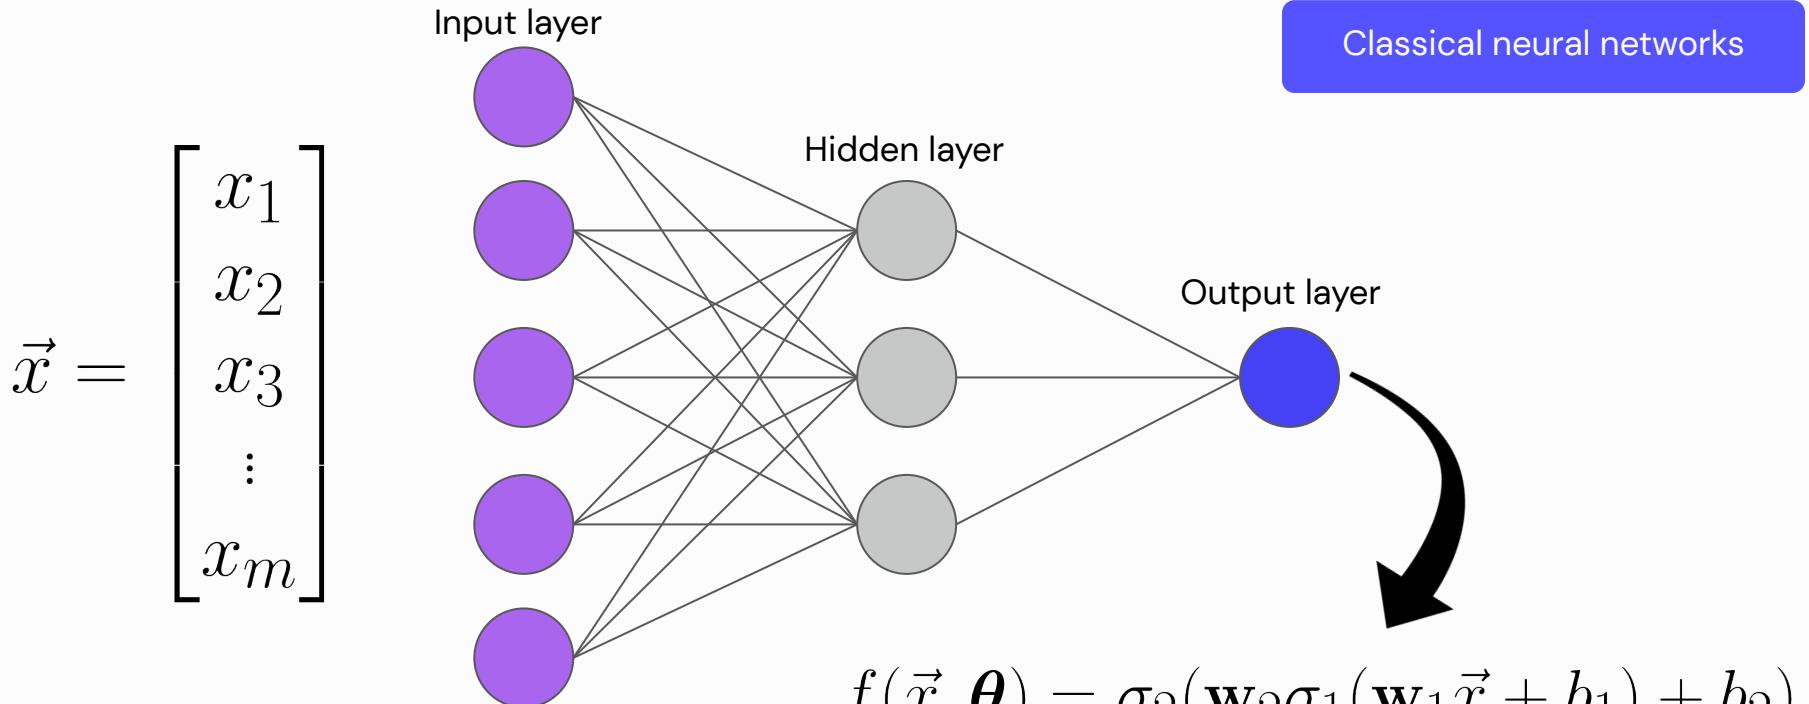
Support vector machines, autoencoders....



*Just two of the most common applications



What is Quantum Machine Learning?

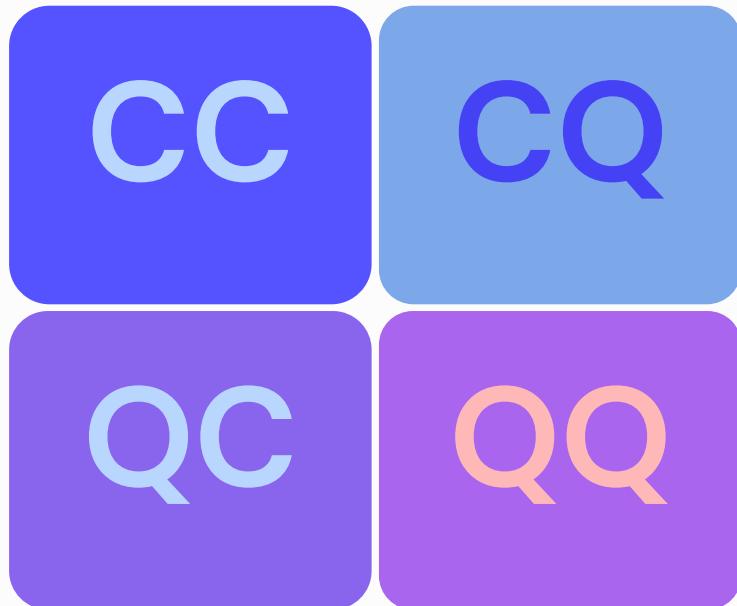




What is Quantum Machine Learning?

Data processing paradigm

Source of data



aQ



What is Quantum Machine Learning?

Data processing paradigm

Source of data

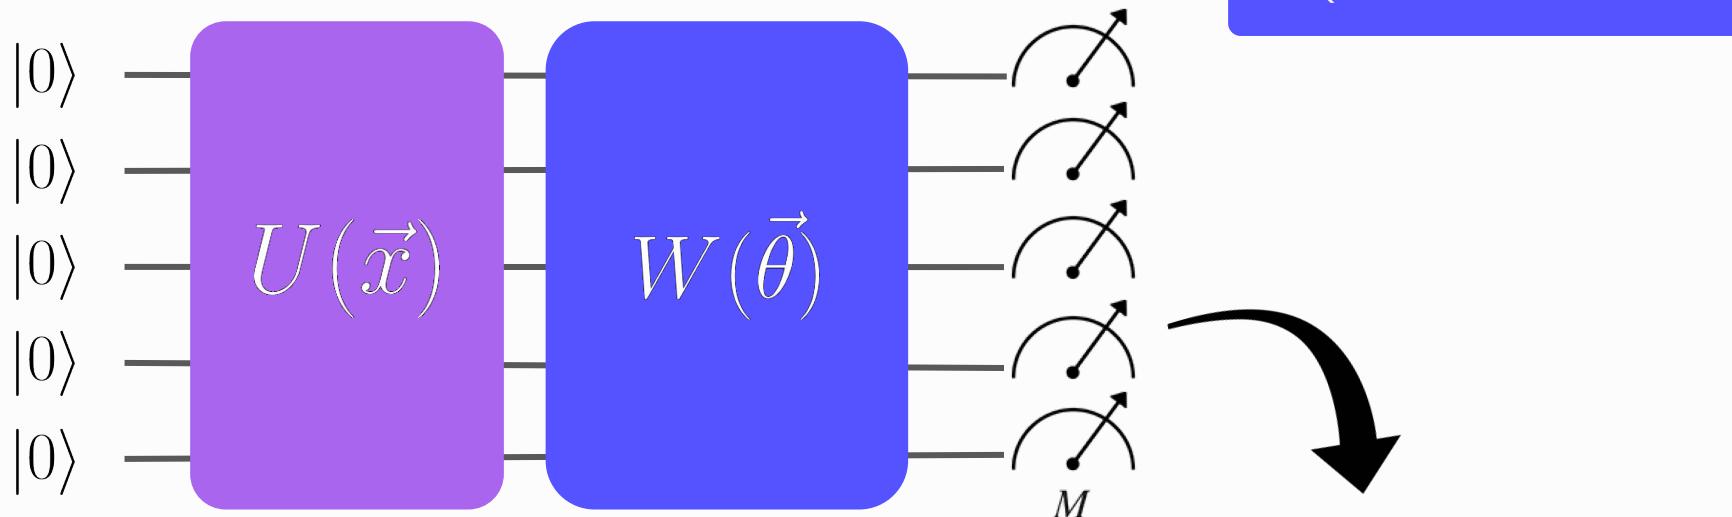


Now popular to work with CQ, but *how does this compare with CC?*

aQ



What is Quantum Machine Learning?



$$f(\vec{x}, \vec{\theta}) = \langle 0 \dots | U^\dagger(\vec{x}) W^\dagger(\vec{\theta}) H W(\vec{\theta}) U(\vec{x}) | 0 \dots \rangle$$



Classical and Quantum Similarity Learning



What is similarity learning?



Similarity learning is the act of learning the notion of what makes two (or more) object similar



One way of handling this classically is to train a **Siamese network**. Basically two neural networks.



To handle this quantumly, we can use quantum similarity networks, or, GQSim

$$f(\vec{x}_i, \vec{x}_j, \vec{\theta}) \approx y_{ij}$$

\vec{x}_i

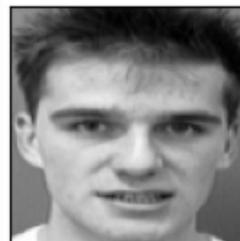


\vec{x}_j



y_{ij}

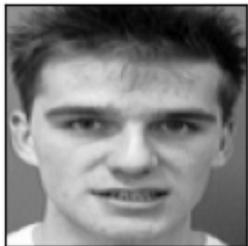
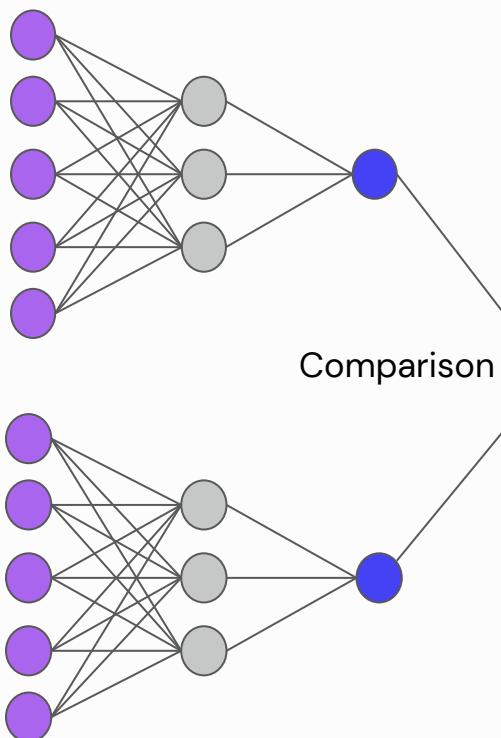
0 (similar)



1 (dissimilar)



Classical Siamese Network

 \vec{x}_i  \vec{x}_j 

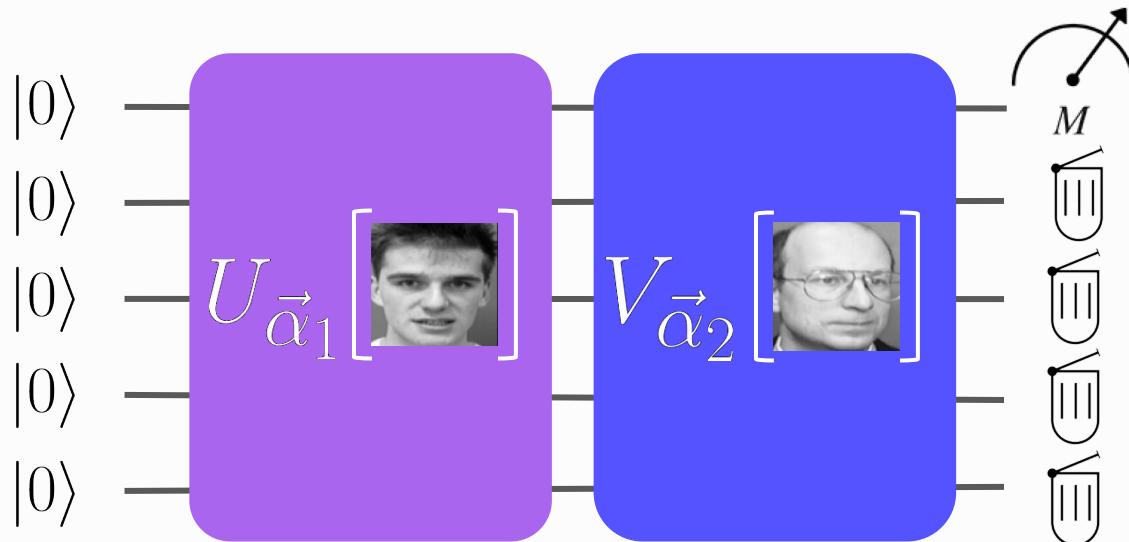
$$f(\vec{x}_i, \vec{x}_j, \vec{\theta}) \approx y_{ij}$$

Activation
 y_{ij}

 PyTorch



Quantum Similarity Network



$$f(\vec{x}_i, \vec{x}_j, \vec{\theta}) \approx y_{ij}$$

Can be as simple as the probability of measuring a single qubit in the 1 or 0 state.

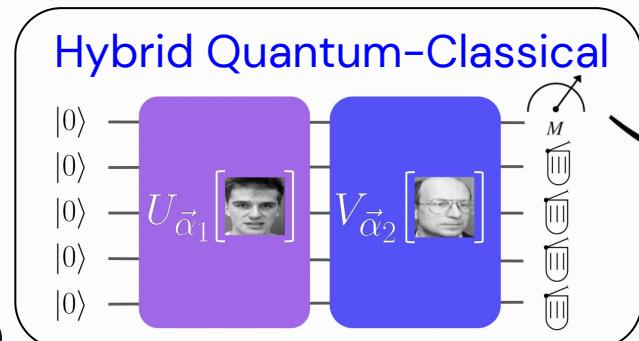
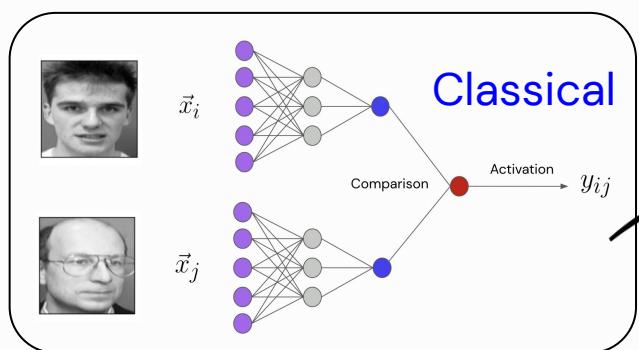


Radha, S.K. and Jao, C., 2022. Generalized quantum similarity learning. *arXiv preprint arXiv:2201.02310*.





Enter Covalent



aQ

Overview

Started - Ended

Sep 16, 15:03:40 - Sep 16, 15:04:59

Runtime

1m

Directory

/Users/jbaker/C...c2-b3e9-4ff9e320957f

Input

```
{"args": [], "kwargs": {"train_path": "/Users/jb...
```

Result

```
(80.0, 82.0, SiameseNetwork(\n    (cnn): Sequent...
```

Executor: task

```
#ct.lattice
def quantum_classical_workflow(train_path, test

    # Run the classical sublattice
    results_classical, loss_history_classical,
    classical_workflow(train_path=pixel_tr...
        test_path=pixel_tes...
        train_batch=train_b...
        pytorch_optimizer=p...
        epochs=epochs,
        lr=lr_c,
        print_intermediate=...
```



Modularize your experiments.

@ct.electron

Define a quantum similarity circuit



@ct.electron

Construct a classical Siamese network

@ct.electron

Construct the contrastive loss function

@ct.electron

Use classical techniques to optimize the network parameters

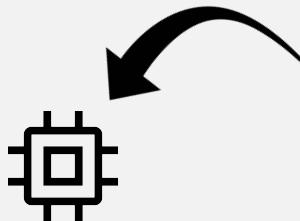
@ct.electron

Test/validate the learned models using new datasets



Covalent.

Modularize your experiments.



Can dispatch specific `@ct.electrons` to *different* remote devices

`@ct.electron`

Construct a classical Siamese network

`@ct.electron`

Define a quantum similarity circuit

`@ct.electron`

Construct the contrastive loss function

`@ct.electron`

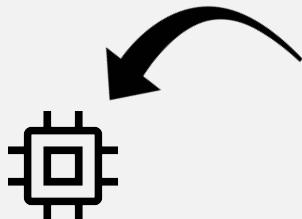
Use classical techniques to optimize the network parameters

`@ct.electron`

Test/validate the learned models using new datasets



Modularize your experiments.



Can dispatch specific `@ct.electrons` to *different* remote devices

`@ct.lattice`

Call all electrons in the desired order to create a **heterogeneous** workflow

`@ct.electron`

Construct a **classical Siamese network**

`@ct.electron`

Define a quantum similarity circuit

`@ct.electron`

Construct the contrastive loss function

`@ct.electron`

Use classical techniques to optimize the network parameters

`@ct.electron`

Test/validate the learned models using new datasets



The workshop



- 1 hour workshop.
- Use classical and quantum methods for similarity learning and compare them.
- Let us know about your own workflows!

similarity_learning.ipynb

Similarity learning with Covalent

In this workshop, you will learn how to use [covalent](#) to manage hybrid quantum-classical workflows for the task of similarity learning. Specifically, within the same workflow, we will dispatch a hybrid quantum-classical machine learning (ML) algorithm and a purely classical ML algorithm to recognize *similar* 2x2 pixel images. Later, we will compare the accuracy to the two approaches with a simple test and finish the workshop investigating the *learnt notion of similarity* in the quantum and classical models. Although we study a toy problem, in the present state of the field of Quantum Machine Learning (QML), the workflows introduced here are exemplar real experiments into the effectiveness of similarity learning with QML.

In the process of doing so, you will be practically introduced the the basic features of [covalent](#) as well as receiving guidance on how to conduct more complex work flows using remote executors and other plugins. This workshop can then serve as guide to *covalentify* your own workflows, with all the advantages that come with it.

For the structure of this notebook, please see the table of contents below:

Table of contents

1. Getting started with Covalent
2. Loading and preparing the training and testing datasets
3. A quick visual introduction to the datasets
4. Classical Siamese networks
5. Quantum similarity networks
6. A heterogeneous workflow: comparing accuracy scores
7. Investigating learnt notions of similarity in the classical and quantum networks
8. Conclusions
9. References

SETUP



- Clone the Github repository

```
git clone https://github.com/AgnostiqHQ/tutorials_coalent_ieee_2022.git
```

- Download [Miniconda](#)
- Install the conda environment (`environment.yml` in root directory)

```
conda env create -f environment.yml
```

- Activate environment

```
conda activate ieee_coalent
```

- Add to kernel list

```
python -m ipykernel install --user --name=ieee_coalent
```